



Self-Adaptive Fine-grained Multi-modal Data Augmentation for Semi-supervised Multi-modal Coreference Resolution

Li Zheng*

zhengli@whu.edu.cn

School of Cyber Science and
Engineering, Wuhan University
Wuhan, China

Boyu Chen*

chenboyu1225@whu.edu.cn

School of Cyber Science and
Engineering, Wuhan University
Wuhan, China

Hao Fei

haofei37@nus.edu.sg

School of Computing, National
University of Singapore
Singapore, Singapore

Fei Li*[†]

foxf823@gmail.com

School of Cyber Science and
Engineering, Wuhan University
Wuhan, China

Shengqiong Wu

swu@u.nus.edu

School of Computing, National
University of Singapore
Singapore, Singapore

Lizi Liao

lzliao@smu.edu.sg

SCIS, Singapore Management
University
Singapore, Singapore

Donghong Ji*

dhji@whu.edu.cn

School of Cyber Science and
Engineering, Wuhan University
Wuhan, China

Chong Teng*

tengchong@whu.edu.cn

School of Cyber Science and
Engineering, Wuhan University
Wuhan, China

Abstract

Coreference resolution, an essential task in natural language processing, is particularly challenging in multi-modal scenarios where data comes in various forms and modalities. Despite advancements, limitations due to scarce labeled data and underleveraged unlabeled data persist. We address these issues with a self-adaptive fine-grained multi-modal data augmentation framework for semi-supervised MCR, focusing on enriching training data from labeled datasets and tapping into the untapped potential of unlabeled data. Regarding the former issue, we first leverage text coreference resolution datasets and diffusion models, to perform fine-grained text-to-image generation with aligned text entities and image bounding boxes. We then introduce a self-adaptive selection strategy, meticulously curating the augmented data to enhance the diversity and volume of the training set without compromising its quality. For the latter issue, we design a self-adaptive threshold strategy that dynamically adjusts the confidence threshold based on the model's learning status and performance, enabling effective utilization of valuable information from unlabeled data. Additionally, we incorporate a distance smoothing term, which smooths distances between positive and negative samples, enhancing discriminative power of

the model's feature representations and addressing noise and uncertainty in the unlabeled data. Our experiments on the widely-used CIN dataset show that our framework significantly outperforms state-of-the-art baselines by at least 9.57% on MUC F1 score and 4.92% on CoNLL F1 score. Remarkably, against weakly-supervised baselines, our framework achieves a staggering 22.24% enhancement in MUC F1 score. These results, underpinned by in-depth analyses, underscore the effectiveness and potential of our approach for advancing MCR tasks.

CCS Concepts

• Computing methodologies → Artificial intelligence.

Keywords

Coreference Resolution, Multi-modal, Semi-supervised Learning

ACM Reference Format:

Li Zheng, Boyu Chen, Hao Fei, Fei Li, Shengqiong Wu, Lizi Liao, Donghong Ji, and Chong Teng. 2024. Self-Adaptive Fine-grained Multi-modal Data Augmentation for Semi-supervised Multi-modal Coreference Resolution. In *Proceedings of the 32nd ACM International Conference on Multimedia (MM '24)*, October 28–November 1, 2024, Melbourne, VIC, Australia. ACM, New York, NY, USA, 10 pages. <https://doi.org/10.1145/3664647.3680966>

1 Introduction

Coreference resolution (CR), which identifies all mentions referring to the same entity, plays a crucial role in many downstream applications, such as relation extraction [11, 49], question answering [5, 46, 55], and sentiment analysis [15, 24, 53, 56, 57]. However, in real-world scenarios, particularly in social media, data often comes in various forms and modalities [16, 28, 29, 43, 44], including text, images, rather than pure text. Therefore, some recent advancements have transferred to multimodal coreference resolution [17, 18], where coreference occurs not only between the

*Key Laboratory of Aerospace Information Security and Trusted Computing, Ministry of Education

[†]Fei Li is the corresponding author.


Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

MM '24, October 28–November 1, 2024, Melbourne, VIC, Australia


© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 979-8-4007-0686-8/24/10

<https://doi.org/10.1145/3664647.3680966>



(a) an example from the CIN dataset



(b) an example of fine-grained multimodal data generation

Figure 1: An example from the CIN dataset (a) [18] and the process of fine-grained multimodal data generation (b).

entities in text but also the objects in image, as shown in Figure 1 (a). Compared to the text-based coreference resolution, Multimodal Coreference Resolution (MCR) presents even greater challenges due to the substantial semantic gap between different modalities and the scarcity of annotated data.

Existing works have made commendable efforts in MCR. Goel et al. [18] proposed a multimodal pipeline and utilized weak supervision to identify coreference chains. In contrast, another study by Goel et al. [17] argued that weak supervision alone is unable to resolve the ambiguity of multiple instances of the same object class. To address this issue, they introduced semi-supervised learning and employed cross-modal attention. Despite these advances, current methods in MCR still tied to model optimization and extracting information features. Data, as a key dimension that can significantly affect model performance, has been far less explored. Therefore, we summarize the challenges facing current research from a data-centric perspective in two aspects:

- **Scarcity of labeled data.** The manual annotation of large-scale fine-grained multimodal datasets is costly, resulting in a scarcity of extensively labeled data required for research on multimodal coreference resolution. This scarcity poses significant challenges in training MCR models. Currently, automatic data expansion [8, 23] is widely used as a practical technique to alleviate the data scarcity problem. Xiao et al. [47] applied the stable diffusion [35] to generate high-quality images with a given text input to solve the image captioning task, but the method is limited to coarse-grained expansion of text-image pairs, ignoring fine-grained objects, entities, etc. For the coreference resolution task, existing methods [10] mainly focus on textual coreference resolution data augmentation, while the MCR data augmentation remains largely unexplored. Therefore, how to adaptively expand high-quality fine-grained multimodal labeled data is the essence of successful MCR.

- **Under-exploitation of unlabeled data.** On the other hand, despite addressing the fundamental issue of scarcity of labeled data through adaptive data augmentation, the untapped potential of unlabeled data remains underexplored. In semi-supervised learning,

existing methods [17] set a fixed threshold to ensure the quality of pseudo-labels for unlabeled data. Unfortunately, a fixed high threshold (e.g., 0.9) poses several limitations: (1) Loss of useful information: By using a fixed high threshold, many unlabeled data that provide valuable insights are filtered out, potentially discarding useful samples. (2) Limitation on the quantity of training data: With a fixed high threshold, only a limited number of unlabeled data are selected, restricting the scale of training data. (3) Ignoring the model's learning status: A fixed threshold ignore the dynamic nature of training process as it cannot dynamically adjust the threshold based on the model's learning status. Therefore, it is crucial for the MCR task to effectively harness valuable information from unlabeled data based on the learning status of the model in order to train more robust and accurate MCR models.

Taking into consideration the two aforementioned aspects, we propose a novel framework: *Self-adaptive Labeled and Unlabeled multimodal Data Augmentation (SLUDA)*, to improve semi-supervised MCR. **Firstly**, as depicted in Figure 1 (b), we utilize a text CR dataset [32] and a diffusion model [27], as well as randomly generated bounding boxes, to generate text-image pairs and object-entity pairs for fine-grained data expansion. Subsequently, we design a self-adaptive selection strategy to filter and select the generated multimodal data and combine it with existing labeled data for model training. This enables us to increase the quantity and diversity of data while ensuring quality, achieving the first goal of data augmentation for labeled data. **Next**, we adopt a self-adaptive threshold strategy to fulfill the second goal of fully tapping into the unlabeled data. By dynamically adjusting the confidence threshold based on the model's learning status and performance, we estimate the threshold based on the exponential moving average (EMA) of confidence scores from unlabeled data. This allows for the utilization of a lower threshold during the early stage of training, and helps the model to accelerate convergence. As training progresses, the threshold gradually increases to filter out unreliable pseudo-labels and enhance the accuracy of the model. **Furthermore**, considering that pseudo-labeled data are unavoidably noisy than manually-labeled data, we introduce a distance smoothing term. It smooths the distances between positive and negative samples, aiding the model in learning more discriminative feature representations and improving the utilization of unlabeled data.

To verify the effectiveness of our model, we conduct experiments on the benchmark CIN dataset [18]. The results demonstrate that our framework significantly helps the state-of-the-art model to further improve its performance by at least 9.57% on MUC F1 score and 4.92% on CoNLL F1 score. Most strikingly, compared to the weakly-supervised baselines, ours outperforms their MUC F1 score by 22.24%. Additionally, we conduct extensive experiments to demonstrate the effectiveness of adaptive selection and utilization of labeled and unlabeled data in our framework. Our main contributions are summarized as follows:

- We propose self-adaptive expansion and selection techniques to enrich the training data with generated high-quality MCR data, effectively addressing the issue of labeled data scarcity.
- We devise a self-adaptive threshold strategy to leverage unlabeled data, achieving a balance between quality and quantity. Moreover, we introduce a distance smoothing policy to enhance the discriminative ability of feature representations.

- Our extensive experimental results on the CIN dataset demonstrate that our SLUDA framework achieves the SoTA performance and outperforms the best baseline with large margins.

2 Related works

2.1 Multimodal Coreference Resolution

Within the multimodal learning community [12–14, 26, 42, 45, 54], there has been growing interest in multimodal coreference resolution. Several studies [17–19] have been conducted on this topic. Goel et al. [18] introduced the CIN dataset, which addresses the problem of coreference resolution in long narratives within visual scenes. They proposed a multimodal pipeline and utilized weak supervision to learn to identify coreference chains. Goel et al. [17] proposed a semi-supervised training approach based on a fixed threshold to effectively learn from unlabeled sets. However, despite achieving partial success, these methods suffered the issue of scarce labeled data and failed to fully leverage the potential of unlabeled data.

2.2 Semi-supervised Learning

Semi-supervised learning is an effective paradigm that utilizes a large amount of unlabeled data along with limited labeled data. Various methods [9] have been proposed in the field of semi-supervised learning, such as pseudo-labeling and consistency regularization. In pseudo-labeling, the models [34, 48] use unlabeled samples with high confidence as training targets, thereby reducing the density of data points at the decision boundary. In consistency regularization, the models [2, 4] are self-supervised on unlabeled data, providing additional training signals. To mitigate the confirmation bias in pseudo-labeling, Goel et al. [17] proposed a threshold-based technique to ensure the quality of pseudo-labels, where only unlabeled data with confidence above the threshold are retained. While promising results have been achieved, a fixed high threshold can lead to the disregard of a large number of ambiguously predicted unlabeled examples, especially in the early and mid-training stages.

2.3 Multimodal Data Augmentation

Multimodal data augmentation is the process of creating additional training data with high quality and diversity. It has been widely applied in various machine learning tasks [1, 20, 50, 58]. Recently, Wang et al. [40] proposed a framework for paired cross-modal data augmentation, which generates an infinite amount of paired data to train cross-modal retrieval models. Xiao et al. [47] utilized stable diffusion models for text-to-image generation, expanding the training set with high-quality image-caption pairs. In this paper, we employ a diffusion model [27] to augment fine-grained labeled data, addressing the issue of scarce labeled data, enhancing the diversity and richness of the MCR dataset, and improving the generalization and performance of MCR models.

3 Methodology

In this paper, we propose a Self-adaptive Labeled and Unlabeled multimodal Data Augmentation (SLUDA) framework that fully utilizes data from two perspectives to address the MCR task. The architecture is shown in Figure 2. First, we generate fine-grained image-caption pairs and adaptively select high-quality generated

MCR data for labeled data augmentation. Then, we design a self-adaptive threshold strategy, considering factors such as data quality and model confidence, and introduce a distance smoothing term to enhance the utilization of unlabeled data. Finally, with preprocessed multimodal labeled and unlabeled data, we employ two modality-specific Variational Autoencoders (VAEs) to learn latent representations for different modalities and perform the MCR task.

3.1 Task Definition

Let (I, C) represent an image-caption pair, where C describes the image I as shown in Figure 1 (a). We denote the set of p mentions as $M = \{m_1, m_2, \dots, m_p\}$, and the image I contains q regions denoted as $I = \{r_1, r_2, \dots, r_q\}$. The objective of multimodal coreference resolution is to identify text mentions that refer to the same entity and associate each mention with a specific region in the image.

3.2 Labeled Data Augmentation

In the task of multimodal coreference resolution, a large number of fine-grained annotated image-caption pairs are typically required. Existing MCR datasets, such as CIN [18], demand human annotators to label entity mentions referring to the regions in the referenced image, construct coreference entity chains, and draw bounding boxes in the image. This annotation process is not only labor-intensive but also time-consuming. Furthermore, the collected images and annotated mentions may suffer from incompleteness and lack diversity, which limits the generalization capabilities of models trained on such datasets. To address the challenge of scarce labeled data, we propose a novel approach that leverages existing text-based coreference resolution datasets in conjunction with a diffusion model to generate images, and then adaptively select high-quality multimodal data. By doing so, we mitigate the limitations posed by data scarcity in the MCR task.

Labeled Data Generation. The diffusion models [36, 47, 51] have been proposed and applied for data augmentation, which is a generative model that generates diverse samples by iteratively diffusing noise signals. For the task of grounded text-to-image generation, Li et al. [27] introduced GLIGEN, which builds upon existing pre-trained text-to-image diffusion models and extends their capabilities to incorporate grounding inputs as conditions. We employ GLIGEN and leverage the widely-used text coreference resolution dataset English OntoNotes 5.0 [32] as the foundational dataset for data augmentation. Given the text as the caption $C = \{w_1, w_2, \dots, w_n\}$, a set of mentions $M = \{m_1, m_2, \dots, m_p\}$, and their randomly generated corresponding bounding boxes $B = \{b_1, b_2, \dots, b_q\}$ as input, we aim to generate the image I :

$$I = \text{GLIGEN}(C, M, B), \quad (1)$$

where $I = \{r_1, r_2, \dots, r_q\}$ represents the image I with q regions. We discard images that contain NSFW (Not Safe for Work) content. Then, we pair the generated images with the corresponding caption to form an augmented labeled dataset.

Self-adaptive Selection. Due to the uncertain quality of the generated images, we employ a self-adaptive selection strategy to filter the data based on the quality assessment of the image-caption pairs, ensuring both the quality and quantity of the training set. Specifically, since CLIPScore [21] is used to assess whether the

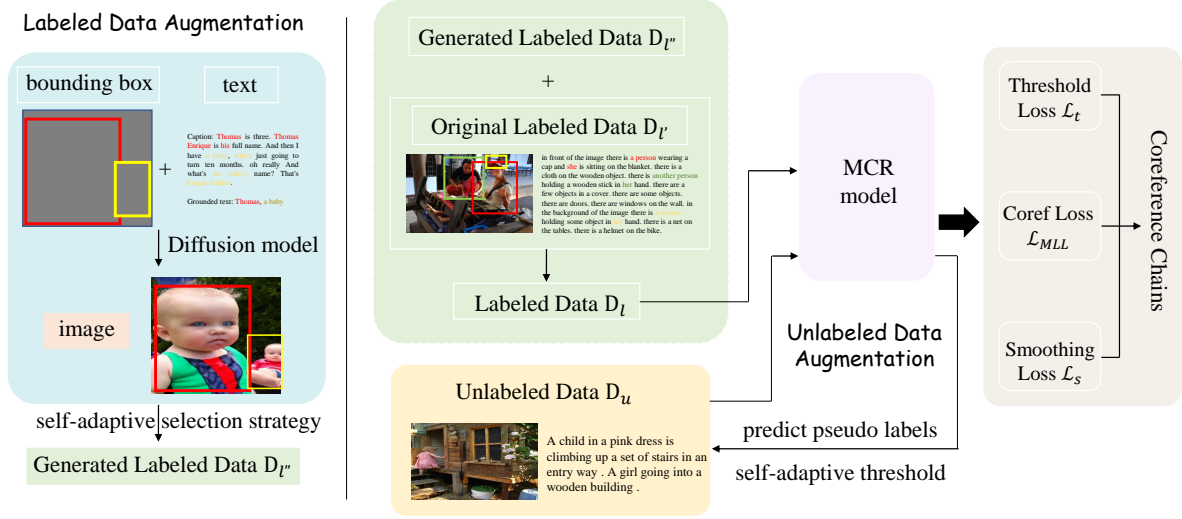


Figure 2: The overall architecture of our model. Initially, we employ a diffusion model to expand the fine-grained MCR labeled data and design a self-adaptive selection strategy to filter and obtain the generated labeled data $D_{L''}$. Subsequently, $D_{L''}$ is merged with the original labeled data D_L , and they are combined with the unlabeled data D_u as input to the MCR model. Then, a self-adaptive threshold strategy is adopted to fully tap into the valuable information of unlabeled data.

generated image can match and correlate with the given textual description, we evaluate the generated image-caption pairs using CLIPScore and sort them according to their evaluation scores. During the initial stages of training, we incorporate all the generated data as labeled data into the training set to ensure sample diversity. As training progresses, we filter out a portion of low-scoring generated data based on the model's learning status to maintain the quality of the labeled data. The learning status of the model is estimated using the predicted scores, which are then smoothed using Exponential Moving Average (EMA) [31]. EMA calculates the current score as a weighted average of the current and previous EMA values, providing a more stable reflection of the model's learning status. The filtering ratio τ_t is defined as follows:

$$EMA = \alpha_1 * o_{r_t} + (1 - \alpha_1) * \tau_{t-1}, \tau_t = \exp(-\beta * EMA), \quad (2)$$

where α_1 is the smoothing factor ($0 \sim 1$), o_{r_t} represents the prediction accuracy at time t , β serves as a control parameter for the filtering intensity. When the EMA value is higher, indicating better overall predicted scores and superior performance of the model in handling the task, we filter out a larger proportion of generated data to ensure the quality of labeled data. Conversely, when the model's learning status is poorer, we filter out a smaller amount of generated data, allowing for the retention of a larger number of samples for training. This approach enables a balance between the quality and quantity of the training set throughout the training process, effectively leveraging the labeled data to enhance the model's performance. This generated and adaptively filtered labeled dataset is denoted as $D_{L''}$, which is merged with the labeled dataset D_L from the multimodal dataset CIN. Therefore, the final labeled dataset is $D_L = D_{L''} + D_L$.

3.3 Unlabeled Data Augmentation

From the perspective of unlabeled data, existing method [17] fails to fully utilize the potential of unlabeled data, resulting in the waste

of unlabeled data and a decrease in the robustness of MCR task. Moreover, unlabeled data often contains noise, and reducing the impact of noise on the model while promoting more accurate and reliable feature representation learning is crucial for MCR. Thus, we propose self-adaptive threshold and distance smoothing techniques to effectively leverage unlabeled data and mitigate noise within it.

Self-adaptive Threshold. Considering that the fixed high threshold method results in the wastage of numerous samples, the samples with ambiguous predictions contribute negligibly during the training phase. Hence, to harness unlabeled data effectively, we propose a self-adaptive threshold strategy that reflects the learning status of the model. The strategy estimates the effectiveness of learning based on the model's prediction confidence and subsequently calculates the threshold τ_t using the exponential moving average of the confidence. The threshold τ_t is defined and adjusted as follows:

$$\tau_t = \alpha_2 * o_t + (1 - \alpha_2) * \tau_{t-1}, \quad (3)$$

where t represents the time step of training, o_t represents the prediction accuracy at time t , and α_2 is the smoothing factor ($0 \sim 1$). By employing this strategy, we can use a lower threshold in the early stages to leverage unlabeled data and accelerate the model convergence. As the training progresses, we gradually increase the threshold to filter out unreliable pseudo-labels and improve the model's accuracy. For unlabeled data, we focus on using cross-entropy loss with a confidence threshold for entropy minimization-based pseudo-labeling:

$$\mathcal{L}_t = - \sum_{m \in N} \sum_{r \in I} \mathbb{I}(p(y|m, r)) > \tau_t \log(p(y|m, r)), \quad (4)$$

where $p(\cdot)$ is the output probability from the model and \mathbb{I} is the indicator function for confidence-based thresholding with τ being the threshold.

Distance Smoothing. To address the issue of noise in unlabeled data, we introduce a distance smoothing term to smooth the

distances between positive and negative samples. This approach aims to facilitate the model’s ability to learn more discriminative feature representations and maximize the utilization of unlabeled data. First, we compute the cosine similarity scores between each pair of samples to determine whether they are positive or negative samples. For each mention m , we compare its similarity score with another mention m' and use a predefined threshold to determine their labels. If the similarity score exceeds the threshold, it is labeled as a positive sample. Otherwise, it is labeled as a negative sample. Next, for each pair of positive samples (m_i, m_j) , we calculate their Euclidean distance $d_{ij} = \|m_i - m_j\|$. Similarly, for each pair of negative samples (m_i, m_k) , we calculate their Euclidean distance $d_{ik} = \|m_i - m_k\|$. To smooth the distance differences between positive and negative samples, we introduce a smoothing function $g(d_{ij}, d_{ik})$ defined as:

$$g(d_{ij}, d_{ik}) = 1 / (1 + \exp(-\gamma * (d_{ij} - d_{ik}))) . \quad (5)$$

This smoothing function maps the distances between positive and negative samples to a smooth value between 0 and 1, providing a smoothing effect by reducing the distance disparity. Finally, we incorporate the distance smoothing term into the loss function to encourage the model to learn discriminative feature representations and effectively utilize unlabeled data. The formula for the loss function is as follows:

$$\mathcal{L}_s = \sum_{i=0}^p \left[\frac{1}{|\mathcal{P}|} \sum_{j \in \mathcal{P}} d_{ij}^2 + \frac{1}{|\mathcal{N}|} \sum_{k \in \mathcal{N}} \max(0, \lambda - d_{ik})^2 \right] * g(d_{ij}, d_{ik}) , \quad (6)$$

where \mathcal{P} is the positive set, \mathcal{N} is the negative set. The parameter λ controls the separation between positive and negative samples. We design this loss function to narrow the gap between positive examples and widen the gap between positive and negative examples. The smoothing function $g(d_{ij}, d_{ik})$ has a value close to 1 when the difference between d_{ij} and d_{ik} is large, increasing the difference in distance between positive and negative samples. When the difference between d_{ij} and d_{ik} is small, the value of $g(d_{ij}, d_{ik})$ is close to 0, decreasing the distance difference between positive and negative samples. By employing the distance smoothing method, we are able to smooth the distance disparities between positive and negative samples, reduce the impact of noise on the model, and facilitate more accurate and reliable feature representation learning, thereby enhancing the model’s performance.

3.4 MCR Model

The details of our model are shown in Figure 3. For fair comparison, following Goel et al. [17], we utilize BERT [7] as the text encoder to obtain contextualized word representations and Faster RCNN [33] as the visual encoder to obtain region representations of images. To effectively integrate textual and visual features, we employ two modality-specific Variational Autoencoders (VAEs) to obtain latent representations of the two modalities’ features. Then, by applying Product of Experts (PoE) [22] to the latent representations of both modalities, we obtain the multimodal representation.

Text Feature Extraction. We employ BERT to encode each word in the narrative caption $C = \{w_1, w_2, \dots, w_n\}$, as mentioned in Goel et al. [17]. We prepend a [CLS] token and append a [SEP] token

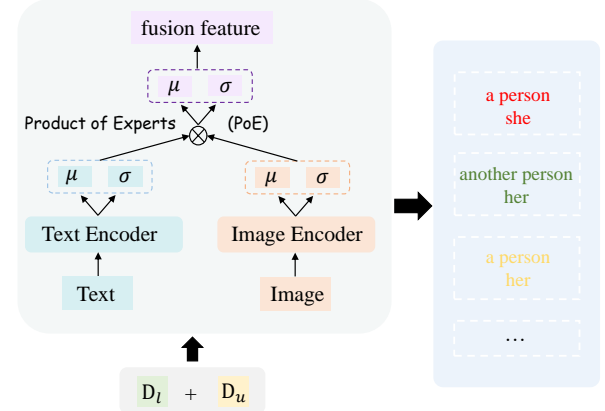


Figure 3: The details of our MCR model.

to each sentence, resulting in $C_0 = \{[CLS], w_1, w_2, \dots, w_n, [SEP]\}$. We then concatenate them together as the input to BERT to generate contextualized token representations. The mention embeddings $M = \{m_1, m_2, \dots, m_p\}$ are computed by averaging the corresponding word embeddings.

Image Feature Extraction. Following Goel et al. [17], we employ Faster RCNN to extract region representations of the image $I = \{r_1, r_2, \dots, r_q\}$. Each region r_i is represented by a d -dimensional joint embedding $r_i \in \mathbb{R}^d$, which encompasses its visual, semantic, and spatial features.

Multimodal Feature Fusion. Previous research [17] has treated text and image features equally, mapping the concatenated features of both modalities to the same latent representation. However, there exist mismatches between text and image, introducing noise into the models used for prediction. To address this, we utilize modality-specific VAEs to map the features of the two modalities into their respective latent representations with internally correlated distributions. This enables the model to adaptively learn its specific data representation for each modality, better capturing the intrinsic relationships and structures among the features of each modality, helping to address the mismatch between text and image. The encoder of each VAE includes dense layers that map the input features to mean vectors μ and standard deviation vectors σ .

For the text modality, the text feature c is inputted into the encoder of the text-VAE to parameterize the mean vector μ_c and the standard deviation vector σ_c . We can approximate the true posterior distribution $p(z_c|c)$ using these parameters, where the distribution of z_c can be represented as $z_c \sim q(z_c|c) = N(\mu_c, \sigma_c^2)$. For the visual modality, the image feature r is also inputted into the encoder of the image-VAE to obtain an approximation of the true posterior distribution $p(z_r|r)$, where the distribution of z_r can be represented as $z_r \sim q(z_r|r) = N(\mu_r, \sigma_r^2)$. Based on the assumption of conditional independence of multimodal latent representation, the latent distribution $p(z_m|c, r)$ of the multimodal representation can be simplified as the combination of two separate latent distributions $p(z_m|c)$ and $p(z_m|r)$. Thus, we apply the product-of-experts (PoE) to estimate the multimodal latent distribution:

$$p(z_m|c, r) \propto p(z_m|c)p(z_m|r) = q(z_c|c)q(z_r|r) , \quad (7)$$

We assume that the latent representations are independent and follow Gaussian distributions parameterized by mean and standard

Table 1: Coreference resolution results on the CIN dataset [18] from our proposed method and other state-of-the-art multi-modal baselines. * means zero-shot performance and others are semi-supervised performance. In the brackets are the improvements of our model over the best-performing baseline(s).

Method	MUC			B ³			CEAF _e			CoNLL
	R	P	F1	R	P	F1	R	P	F1	F1
• Text-based Methods										
Rule-Based [25]*	5.60	10.13	6.40	/	/	/	/	/	/	/
Neural Coref [25]*	0.11	0.17	0.13	/	/	/	/	/	/	/
longdoc [38]*	7.79	8.43	7.24	62.27	76.10	67.69	48.77	84.95	61.02	45.31
• Multimodal Methods										
VisualBERT [37]*	18.17	6.08	8.06	69.01	36.08	41.03	21.25	57.10	28.67	25.92
UNITER [6]*	16.92	7.15	8.83	68.34	44.29	50.22	28.12	72.78	38.91	32.65
VinVL [52]*	16.76	8.60	9.75	68.49	62.32	61.30	42.88	80.81	53.69	41.58
MAF [41]	19.07	15.62	15.65	/	/	/	/	/	/	/
WS-MCR [18]	24.87	18.34	19.19	/	/	/	/	/	/	/
Semi-MCR [17]	<u>31.11</u>	<u>35.25</u>	<u>31.86</u>	<u>70.63</u>	<u>87.85</u>	<u>78.06</u>	<u>63.99</u>	<u>93.44</u>	<u>75.47</u>	<u>61.79</u>
Ours	39.83 (+8.72%)	44.70 (+9.45%)	41.43 (+9.57%)	72.42 (+1.79%)	90.27 (+2.42%)	80.32 (+2.26%)	66.51 (+2.52%)	96.63 (+3.19%)	78.38 (+2.91%)	66.71 (+4.92%)

Table 2: Comparison of narrative grounding performance on the CIN dataset. In the brackets are the improvements of our model over the best-performing baseline(s).

Method	Noun Phrases	Pronouns	Overall
MAF [41]	21.60	18.31	20.91
WS-MCR [18]	30.27	25.96	29.36
Semi-MCR [17]	<u>32.58</u>	<u>28.45</u>	<u>31.71</u>
Ours	36.84 (+4.26%)	32.14 (+3.69%)	35.72 (+4.01%)

deviation. Therefore, the distribution of z_m can be represented as $z_m \sim N(\mu_m, \sigma_m^2)$, where $\mu_m = \frac{\mu_c \sigma_r^2 + \mu_r \sigma_c^2}{\sigma_c^2 \sigma_r^2}$ and $\sigma_m^2 = (\frac{1}{\sigma_c^2} + \frac{1}{\sigma_r^2})^{-1}$. The latent variable z_m of the multimodal representation can be computed as $z_m = \mu_m + \sigma_m^2 \odot \epsilon$, where $\epsilon \sim N(0, I)$.

Training. After obtaining preprocessed multimodal labeled and unlabeled data, we feed them into our model to learn latent representations of different modalities and perform coreference resolution. We calculate the score for the multimodal mention pair (z, z') as $S_{mcr} = \frac{z \cdot z'}{|z||z'|}$. This score measures the likelihood of coreference between z and z' . A higher score indicates a stronger indication of coreference, whereas a lower score suggests non-coreference. We train the coreference scorer by utilizing the negative log marginal likelihood (\mathcal{L}_{MLL}). Finally, the total loss function is the summation of all the aforementioned losses:

$$\mathcal{L} = \mathcal{L}_{MLL} + \eta_1 \mathcal{L}_s + \eta_2 \mathcal{L}_t. \quad (8)$$

4 EXPERIMENTS

4.1 Experimental Setting

Datasets. We evaluate our approach on the CIN dataset [18], which consists of 1,000 test and 880 validation image-caption pairs. Following Goel et al. [17], we use the annotations from the validation split of the CIN dataset as a small labeled set during training. To augment our data, we utilize the widely used text coreference resolution dataset English OntoNotes 5.0 [32], which contains 3,493 documents. After carefully filtering the documents, we obtain 3,116 multimodal coreference examples as labeled data.

Evaluation Metrics. In terms of evaluation metrics, we align with Goel et al. [17] and use CoNLL F1 for coreference resolution evaluation, which is the average F1 score across three metrics: MUC [39], B³ [3], and CEAF_e [30]. For narrative grounding, we consider a prediction to be correct if the Intersection over Union (IoU) between the predicted bounding box and the ground truth box is greater than 0.5 [18].

4.2 Baseline Systems

To validate the effectiveness of our model, we compare it against the following state-of-the-art baselines:

- **VL-BERT:** Su et al. [37] developed VL-BERT with rich aggregation and alignment functions.
- **UNITER:** Chen et al. [6] introduced UNITER, which uses Optimal Transport to align words to image regions.
- **VinVL:** Zhang et al. [52] developed an object detection model to provide object-centric representations of images.
- **MAF:** Wang et al. [41] leveraged fine-grained visual representations to model phrase-object correlation.
- **WS-MCR:** Goel et al. [18] employed weak supervision to learn to identify coreference chains.
- **Semi-MCR:** Goel et al. [17] proposed a semi-supervised approach and designed a multimodal fusion model for MCR.

4.3 Coreference Resolution Results

We comprehensively compare our ALUDA against text-based coreference and multimodal coreference methods on the CIN dataset. The results, presented in Table 1, highlight the superiority of our method over the state-of-the-art (SoTA) baselines and reveals several key findings. Firstly, compared to traditional text-based coreference methods, multimodal methods consistently demonstrate superior performance by leveraging additional visual features. But without carefully incorporating visual information into the task, such as VisualBERT, UNITER and VinVL models, they only achieve marginal

Table 3: Ablation results. The numbers in the brackets are the decreased values compared with our full model.

	CoNLL	MUC	B ³	CEAF _e
Ours	66.71	41.43	80.32	78.38
w/o LDA	63.59(-3.12)	36.92(-4.51)	78.48(-1.84)	75.39(-2.99)
w/o SAT	64.34(-2.37)	38.21(-3.22)	78.86(-1.46)	75.96(-2.42)
w/o SDT	65.46(-1.25)	39.82(-1.61)	79.54(-0.78)	77.04(-1.34)

improvements over text-based methods. MAF and WS-MCR are two weakly supervised methods, and both of them exhibit significant improvements in MUC scores compared to other single-modal and multimodal baselines. By employing the semi-supervised learning method and carefully tuning the model with a small amount of labeled data and a large amount of pseudo-labeled data, Semi-MCR achieves the current state-of-the-art (SoTA) results. Most notably, our model significantly surpasses the state-of-the-art by a substantial margin, with a CoNLL F1 improvement of 4.92%, an MUC F1 improvement of 9.57%. Moreover, compared to MAF and WS-MCR, our method demonstrates remarkable enhancements, with MUC F1 improvements of 25.78% and 22.24%, respectively.

4.4 Narrative Grounding Results

To evaluate the alignment between image regions and phrases in textual data, we set up a narrative grounding task and conduct a comprehensive comparison between our proposed method existing SoTA approaches, including MAF, WS-MCR, and Semi-MCR. The results in Table 2 demonstrate the superiority of our method. Compared to the current state-of-the-art approach (Semi-MCR), our method achieve a 4.26% and 3.69% improvement in noun phrase and pronoun coreference accuracy, respectively. It is worth noting that our proposed method outperforms the MAF, with a 15.24% and 13.83% increase in noun phrase and pronoun coreference accuracy, respectively. These results highlight the advantages of our method and emphasize the necessity of expanding labeled data and effectively leveraging unlabeled data in scenarios where large-scale annotated data is lacking.

4.5 Ablation Study

We perform ablation experiments to evaluate the contribution of each component in our model. As depicted in Table 3, no variant matches the full model’s performance, highlighting the indispensability of each component. Specifically, the absence of labeled data augmentation (LDA) results in the most significant performance drop. The CoNLL F1 and MUC F1 scores decreased by 3.12% and 4.51%, respectively, indicating the crucial role of expanding high-quality labeled datasets in increasing data diversity and model performance. To validate the necessity of self-adaptive threshold (SAT), we remove it and set the threshold to 0.9. The sharp decrease in performance (2.37% drop in CoNLL F1) demonstrates the non-negligible impact of self-adaptive threshold in the semi-supervised MCR task. Furthermore, removing the smoothing distance term (SDT) lead to a significant performance drop, indicating the importance of it in reducing noise in unlabeled data and helping the model learn more discriminative feature representations, thus improving the utilization of unlabeled data.

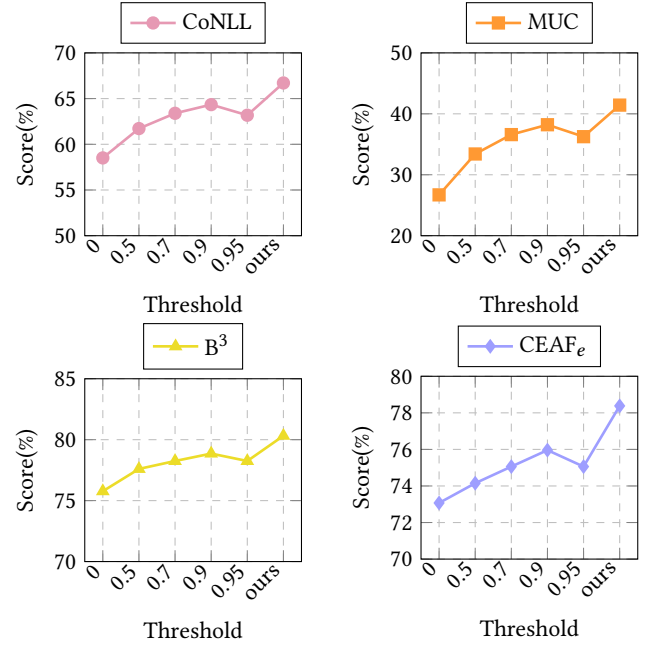


Figure 4: Comparative results between self-adaptive thresholds and different fixed thresholds.

4.6 Deep Analyses on the Effects of Our Proposed Technologies

To further investigate the effectiveness of our method, we conduct in-depth analyses to answer the following questions, with the aim to deeply mine the intuition and analyze implicit phenomena.


Q1: What are the advantages of the self-adaptive threshold?

We validate the effectiveness of our self-adaptive threshold by comparing it with different fixed thresholds in Figure 4. The results indicate that the self-adaptive threshold consistently outperforms the fixed thresholds across all four evaluation metrics. This suggests that the self-adaptive threshold strikes a balance between data quality and quantity, enabling a more flexible selection of unlabeled data. Additionally, we observe that when no filtering of pseudo predictions is applied (i.e., setting the threshold to 0.0), the model performs poorly. This highlights the necessity of a threshold-based training strategy. When consistently using lower fixed threshold, such as 0.5, the model shows some improvement compared to including all pseudo predictions, emphasizing the importance of removing noisy samples to a certain extent. While setting higher fixed threshold, like 0.95, the model’s performance does not improve compared to a threshold of 0.9, in fact, it even declines. This suggests that excessively filtering out unlabeled data result in the loss of high-quality samples, thereby limiting the model’s learning capacity.

Q2: How to evaluate the quality of generated images? We conduct a quality evaluation of the images generated by the diffusion model to quantify and assess the quality of the generated images. To find images that better match the caption and evaluate without relying on real reference images and text, we choose two metrics, MUSIQ and CLIPScore, to evaluate the quality of our generated images. MUSIQ is used to assess the quality of the images themselves, while CLIPScore is used to evaluate whether the generated images

Table 4: Results of the comparison of the self-adaptive selection strategy with different fixed proportion selection strategies.


Training Set	MUC			B^3			CEAF _e			CoNLL
	R	P	F1	R	P	F1	R	P	F1	F1
CIN	37.78	38.65	36.92	72.11	86.62	78.48	63.19	94.92	75.39	63.59
CIN + 25% new	38.56	40.11	38.10	72.21	87.29	78.85	63.78	95.29	75.94	64.29
CIN + 50% new	39.01	40.79	38.72	72.29	87.74	79.09	64.20	95.61	76.36	64.72
CIN + 75% new	<u>39.79</u>	<u>43.75</u>	<u>40.86</u>	72.42	89.46	79.96	65.76	<u>96.35</u>	77.74	<u>66.19</u>
CIN + 100% new	36.42	41.22	37.83	<u>72.71</u>	<u>89.86</u>	<u>80.25</u>	<u>66.60</u>	95.44	<u>77.99</u>	65.36
CIN + self-adaptive	39.83	44.70	41.43	72.42	90.27	80.32	66.51	96.63	78.38	66.71

**Figure 5: Generation image quality evaluation. “Old” means the original data, “All” means the original data and generated data, “New” means the generated data.**

can match and correlate with the given textual descriptions. The evaluation results, as shown in Figure 5, reveal that the generated images (New) achieved high scores on both the MUSIQ and CLIPScore metrics, approaching those of original images (Old). These results indicate that the generated images are of high quality and align well with the given textual descriptions.

Q3: What are the advantages of the self-adaptive selection strategy? To validate the effectiveness of the self-adaptive selection strategy, we compare the results of it with different fixed ratio selection strategies, as shown in Table 4. From the evaluation of image quality, we observe that the overall image quality is high, but there are still a few low-quality images. Thus, we sort the generated images based on their CLIPScore scores and select a portion of them to merge into our training set according to a certain ratio. We find that our self-adaptive selection strategy outperforms all the fixed ratio data selection strategies. We also observe that using all the generated image-caption pairs for the training set does not yield the best results. This might due to the presence of lower-quality images among the generated samples, necessitating the need for filtering based on certain criteria. Additionally, when using lower ratios of generated data, such as 25% and 50%, we notice a slight decrease in overall performance. This suggests that, apart from high-quality samples, the quantity of samples in the training set is also crucial. In conclusion, the above analysis highlight the importance of quality filtering and emphasizes the need to find the optimal balance between high-quality samples and dataset size.

Q4: What is the role of smooth distance term in performance improvement? We visualize the sample features using t-SNE to reveal the effectiveness of smooth distance term. In Figure 6 (a), when employing smooth distance term, the feature differences between positive and negative samples are highly evident. This indicates that the model is encouraged to learn distinct feature representations

**Figure 6: Visualization of the sample features.**

and achieves better discrimination between positive and negative samples. Conversely, in Figure 6 (b), when smooth distance term is removed, the feature distinctions between positive and negative samples are not significant. This limitation hampers the model’s ability to learn discriminative feature representations. This emphasizes the importance of smooth distance term in addressing noise issues in unlabeled data and enhancing model performance.

5 Conclusion

In this paper, we explore two major challenges in the semi-supervised MCR task: scarcity of labeled data and under-exploitation of unlabeled data, and propose a solution called SLUDA. SLUDA expands high-quality MCR training data through a self-adaptive expansion selection strategy, and makes full use of the unlabeled data through a self-adaptive thresholding strategy. Through experiment evaluation, we have found that all of our claimed innovative approaches and hypotheses have been demonstrated to be effective. One of the most interesting findings is that the performance of the model in the semi-supervised MCR task can be significantly improved by adaptively adjusting the threshold value or ratio according to the learning status of the model, both for unlabeled and labeled data. This finding provides insight into our understanding of the role of data quality and quantity in MCR, and offers great promise for advancing further development of this task.

Acknowledgments

This work is supported by the National Key Research and Development Program of China (No. 2022YFB3103602), the National Natural Science Foundation of China (No. 62176187). This research is also supported by the Ministry of Education, Singapore, under its AcRF Tier 2 Funding (Proposal ID: T2EP20123-0052).

References

- [1] Amirhossein Abaskohi, Sascha Rothe, and Yadollah Yaghoobzadeh. 2023. LMCPPF: Paraphrasing-Guided Data Augmentation for Contrastive Prompt-Based Few-Shot Fine-Tuning. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers), ACL 2023*, Anna Rogers, Jordan L. Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, 670–681.
- [2] Abulikemu Abuduweili, Xingjian Li, Humphrey Shi, Cheng-Zhong Xu, and Dejing Dou. 2021. Adaptive Consistency Regularization for Semi-Supervised Transfer Learning. In *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2021*. Computer Vision Foundation / IEEE, 6923–6932.
- [3] Amit Bagga and Breck Baldwin. 1998. Algorithms for scoring coreference chains. In *The first international conference on language resources and evaluation workshop on linguistics coreference*. 563–566.
- [4] Fan Chen, Yiming Zhang, Yaolin Guo, Zhen Liu, and Shiyu Du. 2023. Semi-Supervised Metallographic Image Segmentation via Consistency Regularization and Contrastive Learning. *IEEE Access* 11 (2023), 87398–87408.
- [5] Lei Chen, Bobo Li, Li Zheng, Haining Wang, Zixiang Meng, Runfeng Shi, Hao Fei, Jun Zhou, Fei Li, Chong Teng, et al. 2024. What Factors Influence LLMs' Judgments? A Case Study on Question Answering. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*. 17473–17485.
- [6] Yen-Chun Chen, Linjie Li, Licheng Yu, Ahmed El Kholy, Faisal Ahmed, Zhe Gan, Yu Cheng, and Jingjing Liu. 2020. Uniter: Universal image-text representation learning. In *European conference on computer vision*. Springer, 104–120.
- [7] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2019*, Jill Burstein, Christy Doran, and Thamar Solorio (Eds.). Association for Computational Linguistics, 4171–4186.
- [8] Lisa Dunlap, Alyssa Umino, Han Zhang, Jiezhong Yang, Joseph E. Gonzalez, and Trevor Darrell. 2023. Diversify Your Vision Datasets with Automatic Diffusion-based Augmentation. In *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023*, Alice Oh, Tristan Naumann, Amir Globerson, Kate Saenko, Moritz Hardt, and Sergey Levine (Eds.).
- [9] Yue Fan, Anna Kukleva, Dengxin Dai, and Bernt Schiele. 2023. Revisiting Consistency Regularization for Semi-Supervised Learning. *Int. J. Comput. Vis.* 131, 3 (2023), 626–643.
- [10] Jie Fang and Peifeng Li. 2020. Data Augmentation with Reinforcement Learning for Document-Level Event Coreference Resolution. In *Natural Language Processing and Chinese Computing - 9th CCF International Conference, NLPCC 2020, Zhengzhou, China, October 14-18, 2020, Proceedings, Part I (Lecture Notes in Computer Science, Vol. 12430)*, Xiaodan Zhu, Min Zhang, Yu Hong, and Ruifang He (Eds.). Springer, 751–763.
- [11] Hao Fei, Jingye Li, Shengqiong Wu, Chenliang Li, Donghong Ji, and Fei Li. 2022. Global Inference with Explicit Syntactic and Discourse Structures for Dialogue-Level Relation Extraction. In *Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence, IJCAI 2022*, Luc De Raedt (Ed.). ijcai.org, 4107–4113.
- [12] Hao Fei, Qian Liu, Meishan Zhang, Min Zhang, and Tat-Seng Chua. 2023. Scene Graph as Pivoting: Inference-time Image-free Unsupervised Multimodal Machine Translation with Visual Scene Hallucination. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 5980–5994.
- [13] Hao Fei, Shengqiong Wu, Wei Ji, Hanwang Zhang, and Tat-Seng Chua. 2024. Dysen-VDM: Empowering Dynamics-aware Text-to-Video Diffusion with LLMs. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 7641–7653.
- [14] Hao Fei, Shengqiong Wu, Wei Ji, Hanwang Zhang, Meishan Zhang, Mong-Li Lee, and Wynne Hsu. 2024. Video-of-thought: Step-by-step video reasoning from perception to cognition. In *Proceedings of the International Conference on Machine Learning*.
- [15] Hao Fei, Shengqiong Wu, Yafeng Ren, and Meishan Zhang. 2022. Matching structure for dual learning. In *Proceedings of the International Conference on Machine Learning*. 6373–6391.
- [16] Hao Fei, Shengqiong Wu, Hanwang Zhang, Tat-Seng Chua, and Shuicheng Yan. 2024. VITRON: A Unified Pixel-level Vision LLM for Understanding, Generating, Segmenting, Editing. (2024).
- [17] Arushi Goel, Basura Fernando, Frank Keller, and Hakan Bilen. 2023. Semi-supervised multimodal coreference resolution in image narrations. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, 11067–11081.
- [18] Arushi Goel, Basura Fernando, Frank Keller, and Hakan Bilen. 2023. Who are you referring to? Coreference resolution in image narrations. In *IEEE/CVF International Conference on Computer Vision, ICCV 2023*. IEEE, 15201–15212.
- [19] Danfeng Guo, Arpit Gupta, Sanchit Agarwal, Jiun-Yu Kao, Shuyang Gao, Arijit Biswas, Chien-Wei Lin, Tagyoung Chung, and Mohit Bansal. 2022. GRAVL-BERT: Graphical Visual-Linguistic Representations for Multimodal Coreference Resolution. In *Proceedings of the 29th International Conference on Computational Linguistics, COLING 2022*, Nicoletta Calzolari, Chu-Ren Huang, Hansaem Kim, James Pustejovsky, Leo Wanner, Key-Sun Choi, Pum-Mo Ryu, Hsin-Hsi Chen, Lucia Donatelli, Heng Ji, Sadao Kurohashi, Patrizia Paggio, Nianwen Xue, Seokhwan Kim, Younggyun Hahm, Zhong He, Tony Kyungil Lee, Enrico Santus, Francis Bond, and Seung-Hoon Na (Eds.). International Committee on Computational Linguistics, 285–297.
- [20] Xiaoshuai Hao, Yi Zhu, Srikanth Appalaraju, Aston Zhang, Wanqian Zhang, Bo Li, and Mu Li. 2023. MixGen: A New Multi-Modal Data Augmentation. In *IEEE/CVF Winter Conference on Applications of Computer Vision Workshops, WACV 2023 - Workshops*. 2023. IEEE, 379–389.
- [21] Jack Hessel, Ari Holtzman, Maxwell Forbes, Ronan Le Bras, and Yejin Choi. 2021. CLIPScore: A Reference-free Evaluation Metric for Image Captioning. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, EMNLP 2021*, Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih (Eds.). Association for Computational Linguistics, 7514–7528.
- [22] Geoffrey E. Hinton. 2002. Training Products of Experts by Minimizing Contrastive Divergence. *Neural Comput.* 14, 8 (2002), 1771–1800.
- [23] Fadoua Khmaissia and Hichem Frigui. 2023. Improving Automatic Target Recognition in Low Data Regime using Semi-Supervised Learning and Generative Data Augmentation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2023 - Workshops*. IEEE, 4931–4939.
- [24] M Hari Krishna, K Rahamathulla, and Ali Akbar. 2017. A feature based approach for sentiment analysis using SVM and coreference resolution. In *2017 International conference on inventive communication and computational technologies (ICICCT)*. IEEE, 397–399.
- [25] Kenton Lee, Luheng He, Mike Lewis, and Luke Zettlemoyer. 2017. End-to-end neural coreference resolution. *arXiv preprint arXiv:1707.07045* (2017).
- [26] Bobo Li, Hao Fei, Lizi Liao, Yu Zhao, Chong Teng, Tat-Seng Chua, Donghong Ji, and Fei Li. 2023. Revisiting disentanglement and fusion on modality and context in conversational multimodal emotion recognition. In *Proceedings of the 31st ACM International Conference on Multimedia*. 5923–5934.
- [27] Yuheng Li, Haotian Liu, Qingyang Wu, Fangzhou Mu, Jianwei Yang, Jianfeng Gao, Chunyuan Li, and Yong Jae Lee. 2023. GLIGEN: Open-Set Grounded Text-to-Image Generation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2023*. IEEE, 22511–22521.
- [28] Lizi Liao, Le Hong Long, Zheng Zhang, Minlie Huang, and Tat-Seng Chua. 2021. MMConv: an environment for multimodal conversational search across multiple domains. In *Proceedings of the 44th international ACM SIGIR conference on research and development in information retrieval*. 675–684.
- [29] Lizi Liao, Yunshan Ma, Xiangnan He, Richang Hong, and Tat-seng Chua. 2018. Knowledge-aware multimodal dialogue systems. In *Proceedings of the 26th ACM international conference on Multimedia*. 801–809.
- [30] Xiaoqiang Luo. 2005. On Coreference Resolution Performance Metrics. In *EMNLP*. 25–32.
- [31] Marcus B Perry. 2010. The exponentially weighted moving average. *Wiley Encyclopedia of Operations Research and Management Science* (2010).
- [32] Sameer Pradhan, Alessandro Moschitti, Nianwen Xue, Olga Uryupina, and Yuchen Zhang. 2012. CoNLL-2012 Shared Task: Modeling Multilingual Unrestricted Coreference in OntoNotes. In *Joint Conference on EMNLP and CoNLL - Shared Task*, Sameer Pradhan, Alessandro Moschitti, and Nianwen Xue (Eds.). Association for Computational Linguistics, 1–40.
- [33] Shaoqing Ren, Kaiming He, Ross B. Girshick, and Jian Sun. 2015. Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks. In *Advances in Neural Information Processing Systems 28: Annual Conference on Neural Information Processing Systems 2015*, Corinna Cortes, Neil D. Lawrence, Daniel D. Lee, Masashi Sugiyama, and Roman Garnett (Eds.). 91–99.
- [34] Mamshad Nayeem Rizve, Kevin Duarte, Yogesh S. Rawat, and Mubarak Shah. 2021. In Defense of Pseudo-Labeling: An Uncertainty-Aware Pseudo-label Selection Framework for Semi-Supervised Learning. In *ICLR*. OpenReview.net.
- [35] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. 2022. High-Resolution Image Synthesis with Latent Diffusion Models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2022*. IEEE, 10674–10685.
- [36] Jordan Shipard, Arnold Wiliem, Kien Nguyen Thanh, Wei Xiang, and Clinton Fookes. 2023. Diversity is Definitely Needed: Improving Model-Agnostic Zero-shot Classification via Stable Diffusion. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2023 - Workshops*. IEEE, 769–778.
- [37] Weijie Su, Xizhou Zhu, Yue Cao, Bin Li, Lewei Lu, Furu Wei, and Jifeng Dai. 2019. Vi-bert: Pre-training of generic visual-linguistic representations. *arXiv preprint arXiv:1908.08530* (2019).

- [38] Shubham Toshniwal, Patrick Xia, Sam Wiseman, Karen Livescu, and Kevin Gimpel. 2021. On generalization in coreference resolution. *arXiv preprint arXiv:2109.09667* (2021).
- [39] Marc B. Vilain, John D. Burger, John S. Aberdein, Dennis Connolly, and Lynette Hirschman. 1995. A model-theoretic coreference scoring scheme. In *Proceedings of the 6th Conference on Message Understanding, MUC 1995*. 45–52.
- [40] Hao Wang, Guosheng Lin, Steven C. H. Hoi, and Chunyan Miao. 2022. Paired Cross-Modal Data Augmentation for Fine-Grained Image-to-Text Retrieval. In *MM '22: The 30th ACM International Conference on Multimedia, 2022*, João Magalhães, Alberto Del Bimbo, Shin'ichi Satoh, Nicu Sebe, Xavier Alameda-Pineda, Qin Jin, Vincent Oria, and Laura Toni (Eds.). ACM, 5517–5526.
- [41] Qinxin Wang, Hao Tan, Sheng Shen, Michael W Mahoney, and Zhewei Yao. 2020. Maf: Multimodal alignment framework for weakly-supervised phrase grounding. *arXiv preprint arXiv:2010.05379* (2020).
- [42] Shengqiong Wu, Hao Fei, Wei Ji, and Tat-Seng Chua. 2023. Cross2StrA: Unpaired Cross-lingual Image Captioning with Cross-lingual Cross-modal Structure-pivoted Alignment. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 2593–2608.
- [43] Shengqiong Wu, Hao Fei, Xiangtai Li, Jiayi Ji, Hanwang Zhang, Tat-Seng Chua, and Shuicheng Yan. 2024. Towards Semantic Equivalence of Tokenization in Multimodal LLM. *arXiv preprint arXiv:2406.05127* (2024).
- [44] Shengqiong Wu, Hao Fei, Leigang Qu, Wei Ji, and Tat-Seng Chua. 2024. NExT-GPT: Any-to-Any Multimodal LLM. In *Proceedings of the International Conference on Machine Learning*.
- [45] Shengqiong Wu, Hao Fei, Hanwang Zhang, and Tat-Seng Chua. 2023. Imagine that! abstract-to-intricate text-to-image synthesis with scene graph hallucination diffusion. In *Proceedings of the 37th International Conference on Neural Information Processing Systems*. 79240–79259.
- [46] Wei Wu, Fei Wang, Arianna Yuan, Fei Wu, and Jiwei Li. 2020. CorefQA: Coreference Resolution as Query-based Span Prediction. In *Proceedings of ACL*, Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel R. Tetreault (Eds.). Association for Computational Linguistics, 6953–6963.
- [47] Changrong Xiao, Sean Xin Xu, and Kunpeng Zhang. 2023. Multimodal Data Augmentation for Image Captioning using Diffusion Models. In *Proceedings of the 1st Workshop on Large Generative Models Meet Multimodal Applications, LGM3A 2023*, Zheng Wang, Cheng Long, Shihao Xu, Bingzheng Gan, Wei Shi, Zhao Cao, and Tat-Seng Chua (Eds.). ACM, 23–33.
- [48] Qizhe Xie, Minh-Thang Luong, Eduard H. Hovy, and Quoc V. Le. 2020. Self-Training With Noisy Student Improves ImageNet Classification. In *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2020*. Computer Vision Foundation / IEEE, 10684–10695.
- [49] Zhongxuan Xue, Rongzhen Li, Qizhu Dai, and Zhong Jiang. 2022. Corefdre: Document-level relation extraction with coreference resolution. *arXiv preprint arXiv:2202.10744* (2022).
- [50] Chenchen Ye, Lizi Liao, Suyu Liu, and Tat-Seng Chua. 2022. Reflecting on experiences for response generation. In *Proceedings of the 30th ACM International Conference on Multimedia*. 5265–5273.
- [51] Zebin You, Yong Zhong, Fan Bao, Jiacheng Sun, Chongxuan Li, and Jun Zhu. 2023. Diffusion Models and Semi-Supervised Learners Benefit Mutually with Few Labels. In *NeurIPS*, Alice Oh, Tristan Naumann, Amir Globerson, Kate Saenko, Moritz Hardt, and Sergey Levine (Eds.).
- [52] Pengchuan Zhang, Xiujun Li, Xiaowei Hu, Jianwei Yang, Lei Zhang, Lijuan Wang, Yejin Choi, and Jianfeng Gao. 2021. Vinvl: Revisiting visual representations in vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 5579–5588.
- [53] Zheng Zhang, Lizi Liao, Minlie Huang, Xiaoyan Zhu, and Tat-Seng Chua. 2019. Neural multimodal belief tracker with adaptive attention for dialogue systems. In *The world wide web conference*. 2401–2412.
- [54] Yu Zhao, Hao Fei, Yixin Cao, Bobo Li, Meishan Zhang, Jianguo Wei, Min Zhang, and Tat-Seng Chua. 2023. Constructing holistic spatio-temporal scene graph for video semantic role labeling. In *Proceedings of the 31st ACM International Conference on Multimedia*. 5281–5291.
- [55] Li Zheng, Hao Fei, Fei Li, Bobo Li, Lizi Liao, Donghong Ji, and Chong Teng. 2024. Reverse multi-choice dialogue commonsense inference with graph-of-thought. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 38. 19688–19696.
- [56] Li Zheng, Donghong Ji, Fei Li, Hao Fei, Shengqiong Wu, Jingye Li, Bobo Li, and Chong Teng. 2023. ECQED: emotion-cause quadruple extraction in dialogs. *arXiv preprint arXiv:2306.03969* (2023).
- [57] Li Zheng, Fei Li, Yuyang Chai, Chong Teng, and Donghong Ji. 2023. A Bi-directional Multi-hop Inference Model for Joint Dialog Sentiment Classification and Act Recognition. In *CCF International Conference on Natural Language Processing and Chinese Computing*. Springer, 235–248.
- [58] Jing Zhou, Yanan Zheng, Jie Tang, Li Jian, and Zhilin Yang. 2022. FlipDA: Effective and Robust Data Augmentation for Few-Shot Learning. In *Proceedings of ACL*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, 8646–8665.